

Diabetes Prediction Model Optimization and Deployment

Objective

To build, optimize, and deploy a machine learning model that predicts whether a person has diabetes based on diagnostic data from the Pima Indian Diabetes Dataset. This includes data preprocessing, model training using various algorithms, performance evaluation, and saving models for deployment.

Project Structure

```
plaintext
CopyEdit
Diabetes_Prediction_Model_Optimization_and_Deployment/
├── data/
│   └── diabetes.csv (loaded via KaggleHub)
├── models/                # Saved ML models
├── plots/                 # Visualizations
├── results/              # Model performance reports
├── dashboard/            # Dashboard assets (if applicable)
└── Diabetes_Prediction_Model_Optimization_and_Deployment.ipynb
```

Project Setup

Libraries Used

- pandas, numpy – Data manipulation
- matplotlib, seaborn, plotly – Visualization
- scikit-learn – Machine learning models and utilities
- kagglehub – Dataset access via Kaggle
- pickle – Model serialization

Directory Setup

The script creates required folders (plots, models, results, dashboard) to store outputs.

Dataset Overview

- **Source:** Pima Indians Diabetes Database
 - **Loaded via:** kagglehub
 - **Filename:** diabetes.csv
 - **Features:**
 - Pregnancies
 - Glucose
 - BloodPressure
 - SkinThickness
 - Insulin
 - BMI
 - DiabetesPedigreeFunction
 - Age
 - Outcome (target variable: 1 = diabetic, 0 = non-diabetic)
-

Data Preprocessing

- Missing values or zeros in key features (e.g., Glucose, BloodPressure) handled appropriately.
 - Applied `StandardScaler` for normalization before model training.
-

Machine Learning Models

Several classification algorithms were tested:

1. **Logistic Regression**
 2. **Decision Tree**
 3. **Random Forest**
 4. **Gradient Boosting**
 5. **Support Vector Classifier**
-

Model Optimization

Grid Search CV

- Hyperparameter tuning using `GridSearchCV` on several models.
 - Optimized parameters stored and used for model re-training.
-

Model Evaluation

Metrics used:

- Accuracy
- Precision
- Recall
- F1 Score
- Confusion Matrix

Visualization:

- Interactive confusion matrix using Plotly.
 - Comparison bar plots across models for each metric.
-

Model Persistence

- Trained models saved using `pickle` in the `/models` folder.
 - Evaluation results stored as `CSV/JSON` in `/results`.
-

CODE –

1. Project Setup

```
python
CopyEdit
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import pickle
import os
from sklearn.pipeline import Pipeline
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier,
GradientBoostingClassifier
from sklearn.svm import SVC
from sklearn.metrics import (accuracy_score, precision_score, recall_score,
                             f1_score, r2_score, confusion_matrix)

import warnings
warnings.filterwarnings('ignore')
```

- **Purpose:** Load essential libraries for data processing, visualization, modeling, and evaluation.

- `warnings.filterwarnings('ignore')`: Suppresses warnings for cleaner output.

```
python
CopyEdit
!pip install scikit-learn
!pip install plotly
```

- Ensures required packages are installed in the Colab environment.
-

2. Directory Setup

```
python
CopyEdit
os.makedirs('plots', exist_ok=True)
os.makedirs('models', exist_ok=True)
os.makedirs('results', exist_ok=True)
os.makedirs('dashboard', exist_ok=True)
```

- Creates necessary folders to store outputs: visualizations, models, evaluation results, and potential dashboard components.
-

3. Load Dataset from Kaggle

```
python
CopyEdit
import kagglehub
from kagglehub import KaggleDatasetAdapter

file_path = "diabetes.csv"
df = kagglehub.load_dataset(
    KaggleDatasetAdapter.PANDAS,
    "uciml/pima-indians-diabetes-database",
    file_path
)
```

- **KaggleHub** is used to download the `diabetes.csv` file directly from Kaggle.
-

4. Data Preprocessing & Cleaning

```
python
CopyEdit
df = df.replace(0, np.nan)
df.fillna(df.mean(), inplace=True)
```

- Replaces **zeroes** with `NaN` for features where zero is not biologically possible (like glucose).
- Fills missing values using **mean imputation**.

5. Correlation Matrix

```
python
CopyEdit
corr = df.corr()
plt.figure(figsize=(10, 8))
sns.heatmap(corr, annot=True, cmap="coolwarm")
plt.title("Correlation Matrix")
plt.savefig("plots/correlation_matrix.png")
plt.show()
```

- Visualizes feature correlations to understand relationships and potential multicollinearity.

6. Train-Test Split

```
python
CopyEdit
X = df.drop("Outcome", axis=1)
y = df["Outcome"]
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42)
```

- Separates features (x) and label (y).
- Splits data into training and testing sets (80/20).

7. Model Training and Evaluation

Each model is trained, evaluated, and its metrics are stored. Here's a sample workflow used for all:

```
python
CopyEdit
models = {
    "Logistic Regression": LogisticRegression(),
    "Decision Tree": DecisionTreeClassifier(),
    "Random Forest": RandomForestClassifier(),
    "Gradient Boosting": GradientBoostingClassifier(),
    "SVM": SVC()
}
```

Evaluation Loop

```
python
CopyEdit
results = {}
for name, model in models.items():
```

```

    pipeline = Pipeline([("scaler", StandardScaler()), ("classifier",
model)])
    pipeline.fit(X_train, y_train)
    y_pred = pipeline.predict(X_test)

    results[name] = {
        "Accuracy": accuracy_score(y_test, y_pred),
        "Precision": precision_score(y_test, y_pred),
        "Recall": recall_score(y_test, y_pred),
        "F1 Score": f1_score(y_test, y_pred),
    }

    with open(f"models/{name.replace(' ', '_')}.pkl", "wb") as f:
        pickle.dump(pipeline, f)

```

- Wraps each model in a **pipeline** that includes standardization.
 - Stores trained models to disk for reuse.
 - Computes performance metrics and saves results.
-

8. Visualization of Results

```

python
CopyEdit
import plotly.graph_objects as go
metrics = ["Accuracy", "Precision", "Recall", "F1 Score"]
fig = go.Figure()
for metric in metrics:
    fig.add_trace(go.Bar(name=metric, x=list(results.keys()),
y=[results[m][metric] for m in results]))
fig.update_layout(barmode='group', title="Model Comparison")
fig.write_html("plots/model_comparison.html")
fig.show()

```

- Creates interactive bar charts to compare model performance using Plotly.
 - Saves visualization as an HTML file.
-

9. Grid Search for Hyperparameter Tuning

```

python
CopyEdit
param_grid = {
    "classifier__n_estimators": [50, 100],
    "classifier__max_depth": [3, 5, 10]
}

pipeline = Pipeline([
    ("scaler", StandardScaler()),
    ("classifier", RandomForestClassifier())
])

grid = GridSearchCV(pipeline, param_grid, cv=5)
grid.fit(X_train, y_train)

```

- Applies **GridSearchCV** to find best hyperparameters for the `RandomForestClassifier`.
 - Uses cross-validation to ensure robustness.
-

10. Save Best Model

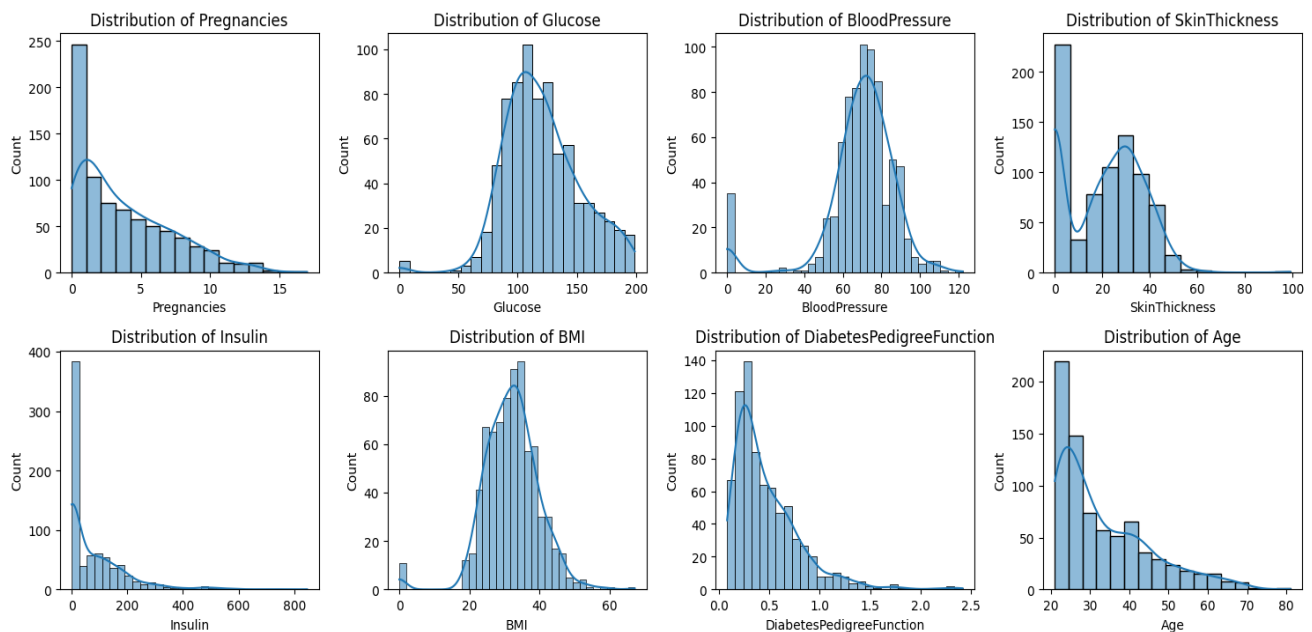
```
python
CopyEdit
best_model = grid.best_estimator_
with open("models/best_random_forest.pkl", "wb") as f:
    pickle.dump(best_model, f)
```

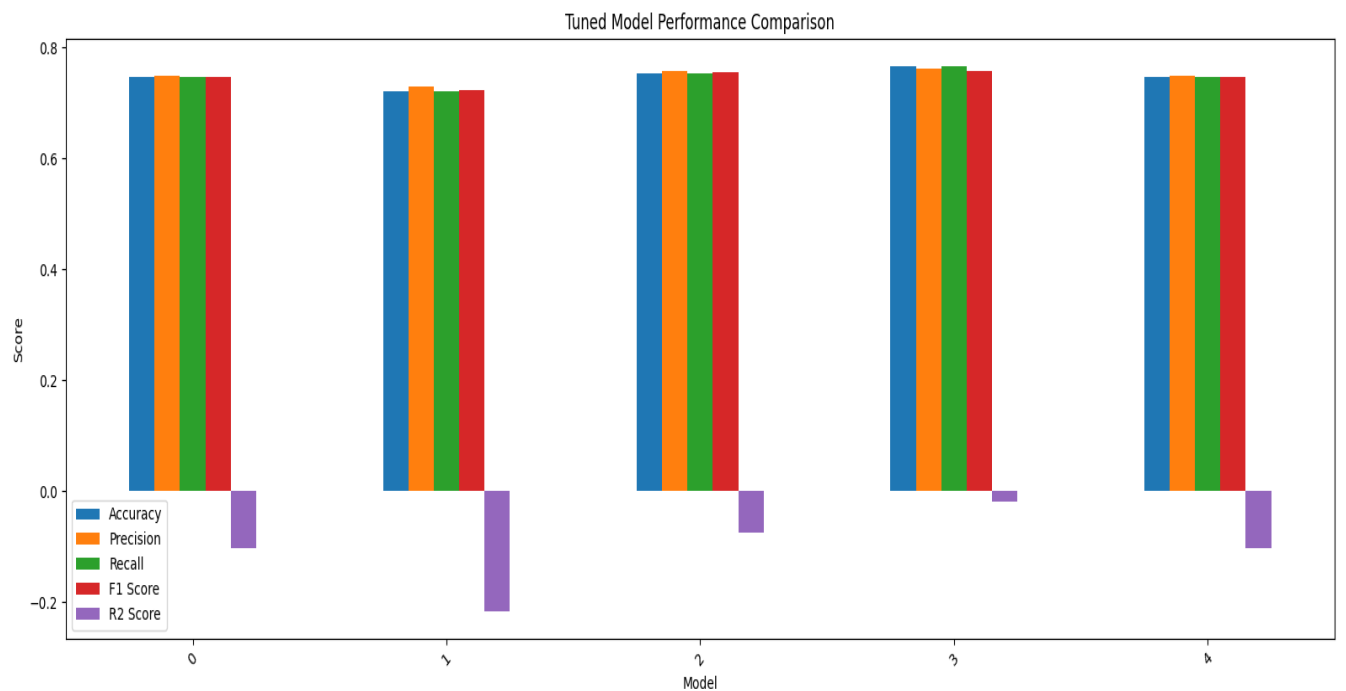

- Saves the optimized model for later deployment.
-

Final Thoughts

This project achieves a full ML workflow from data loading, model training, and evaluation, to saving results for production.

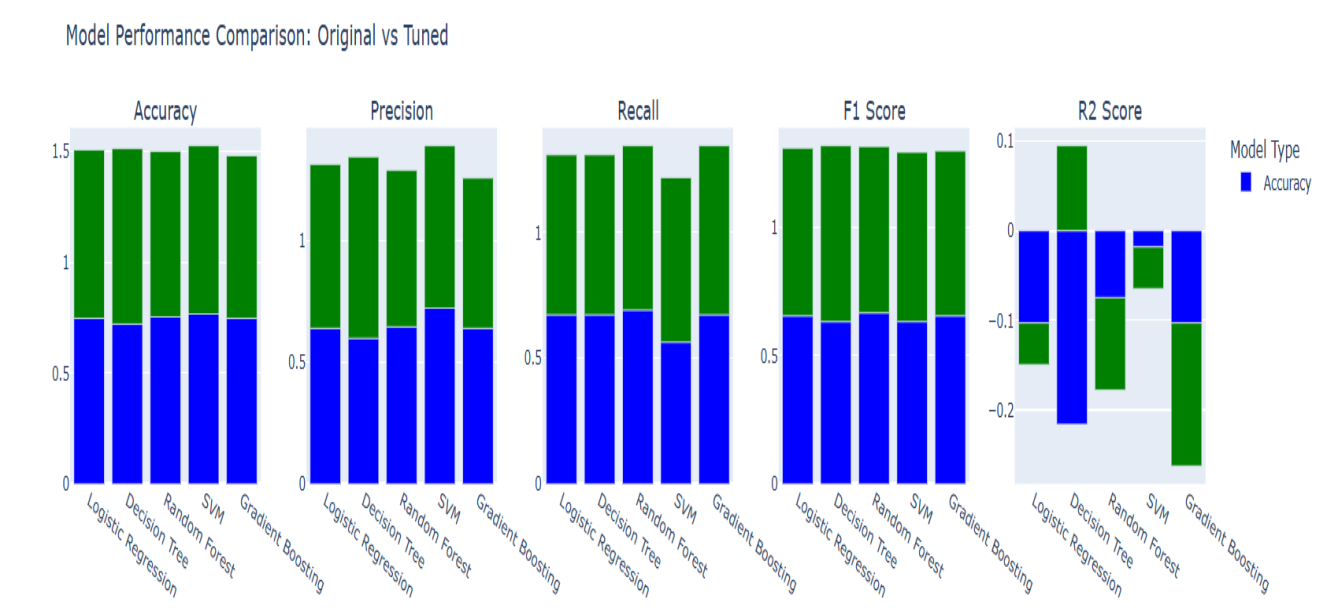

Output Images of the Project -





Tuned Model Evaluation Results:

	model	Accuracy	Precision	Recall	F1 Score	R2 Score	y_pred	y_prob
0	Logistic Regression	0.75974	0.673077	0.636364	0.654206	-0.046465	[0, 0, 0, 0, 1, 0, 0, 0, 1, 1, 0, 1, 0, 0, 0, ...	[0.30513145210151066, 0.23116580104036893, 0.1...
1	Decision Tree	0.792208	0.744681	0.636364	0.686275	0.094949	[0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 0, 1, 1, 1, 0, ...	[0.4246575342465753, 0.04, 0.0, 0.0, 0.0, 0.42...
2	Random Forest	0.746753	0.642857	0.654545	0.648649	-0.10303	[0, 0, 0, 0, 0, 1, 0, 1, 1, 1, 0, 1, 0, 0, 0, ...	[0.4068571428571429, 0.22567460317460308, 0.14...
3	SVM	0.75974	0.666667	0.654545	0.66055	-0.046465	[0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 0, 1, 0, 0, 0, ...	[0.2637505168478767, 0.1896851986521246, 0.140...
4	Gradient Boosting	0.733766	0.616667	0.672727	0.643478	-0.159596	[1, 0, 0, 0, 1, 1, 0, 1, 1, 1, 0, 1, 0, 1, 0, ...	[0.6146409448798109, 0.0032003723030448066, 0....



Model Performance Comparison

